

T S V

ARUN AILESH ANGIPURAPU

Redmond, WA | +1682-392-2299 | tarunsailesh9@gmail.com |

PROFESSIONAL SUMMARY

Data Engineer with 4 years of experience building scalable ETL/ELT pipelines on Azure. Expert in Databricks, Azure Data Factory, and Microsoft Fabric for processing large-scale logistics and business data. Specialized in Medallion Architecture, performance tuning, and ensuring data reliability to provide high-quality, actionable insights for stakeholders.

EXPERIENCE

Quadrant Technologies | Data Engineer

Jan 2024 – Present

- Orchestrated end-to-end ELT pipelines using Azure Data Factory and Databricks to ingest large-scale logistics data into a centralized Medallion Architecture, ensuring a "single source of truth" for global shipping and delivery modules.
- Developed and optimized complex SQL and PySpark transformations to process high-volume logistics events, including Tender Acceptance (TA) and Carrier Planning (CP) metrics, improving data availability for downstream analytics.
- Engineered and maintained robust data models within the Logistics Foundation Model, implementing schema enforcement and data validation to support reliable inventory and transport performance tracking.
- Collaborated with business stakeholders to translate logistics KPIs into interactive Power BI dashboards, utilizing DAX measures to provide real-time visibility into contract compliance and shipping efficiency.
- Streamlined data refresh cycles by tuning Spark jobs and SQL queries, significantly reducing latency for critical logistics reports used by carrier planning teams.
- Implemented automated CI/CD pipelines using Azure DevOps to streamline the deployment of ADF pipelines and Databricks notebooks across Dev, QA, and Prod environments, reducing manual deployment errors by 30%.
- Managed infrastructure as code (IaC) and version control via Git integration in Azure DevOps, ensuring consistent environment configurations and enabling seamless collaboration across the global data engineering team.
- Configured automated unit testing and code quality gates within DevOps pipelines for PySpark scripts, ensuring that only validated, high-quality code was merged into the master branch for production processing.

Wipro Pvt Ltd | Data Engineer

May 2021 – Jan 2023

- Led the end-to-end migration of 60+ Azure Databricks notebooks to Azure Synapse Analytics, re-architecting data pipelines to eliminate premium cluster dependency and achieving a 30% infrastructure cost reduction while maintaining SLA compliance for investor relations workloads.
- Optimized PySpark-based ETL workflows by refactoring expensive transformations, improving partitioning strategies and execution plans, resulting in a 40% reduction in job runtime and enhanced pipeline stability.
- Designed and configured dedicated Apache Spark pools in Azure Synapse, tuning executor memory, dynamic allocation, node sizing, and Spark configuration to deliver scalable, cost-efficient performance aligned with workload characteristics.
- Implemented secured data access patterns using managed identities and controlled mount paths to Azure Data Lake Storage (ADLS Gen2), strengthening data governance, environment isolation (dev/test/prod), and secure cross-environment integration.
- Established CI/CD-driven orchestration and job automation, integrating Git-based deployment workflows with Synapse pipelines and custom PRSML configuration to ensure version-controlled releases, automated validation, and reliable production-grade execution.
- Designed and implemented scalable data models in Synapse to support high-performance analytical queries, enabling faster downstream consumption for investor reporting and executive dashboards.
- Built reusable PySpark-based data transformation frameworks with modularized business logic and metadata-driven configurations, improving maintainability, reducing code duplication across pipelines, and accelerating onboarding of new data sources.

EDUCATION

University of North Texas, Denton, TX

Master of Science in Information Systems

Relevant Coursework: Data Structures & Algorithms, Operating Systems, Computer Architecture, DBMS, Distributed Systems, System Design.

TECHNICAL SKILLS

Programming & Scripting: Python, SQL, PySpark

DataEngineering: AzureDatabricks,AzureSynapseAnalytics,AzureDataFactory,Microsoftfabric,Spark

CloudPlatforms: MicrosoftAzure(DataLakeStorageGen2,BlobStorage,Synapse,PowerApps),AzureDevOps

DataVisualization: PowerBI,Tableau,Excel(PowerPivot,PowerQuery)

Databases: SQLServer,MySQL,AzureSQL,DatabricksLakehouse

VersionControl&CI/CD:Git,GitHub,GitLab,AzureRepos

CERTIFICATIONS

- AzureAIEngineerAssociate(AI-102) Microsoft
- FabricDataEngineerAssociate(DP-700) Microsoft
- DataEngineerProfessional Databricks